Guidelines for Optional Final Project and Proposal

If you choose to do a project, you can extend an idea we've talked about in class or the homework or you can propose your own idea. The project will consist of code implementing your idea and a paper explaining what you did. This is meant to be an individual project.

Your proposal should be around 1 page in length and should include a description of your idea, what you're hoping to accomplish, and suggestions for how you will handle finding data and performing evaluation.

Paper Details:
The paper should be around 5 pages in length. It should include a description of your proposal and the goal you were hoping to accomplish. It should discuss the intuitions behind your approach – why did you expect it to work, was it informed by any linguistic intuitions. The paper should include a description of what you did – what did you implement, how does it work. You should present the results – how does it perform. The paper should also include analysis of your results - why do you think your model performed as it did, were there any surprising results, do you see any areas for improvement, possible future extensions, think about how it could work with data from a different domain or language (think of these as starting points to get you thinking – not as a list of questions to answer).

With your project, you should also submit some technical details:
You should submit your code, instructions for how to run it, data you used, screenshots of it running, and anything else necessary to demonstrate that it works as intended.

When deciding on your idea, think about the following issues that will be important for completing the project:

Data:
What data will you use?
A dataset you've found
Dataset that you will create

Evaluation:
How will you evaluate your performance?
Is there a baseline to compare to?
How is performance on the task usually measured?

Focus on proposing and trying something interesting rather than performance. A project that implements an interesting idea and includes thoughtful analysis even if the results do not show a great improvement is better than a project that achieves great results but with a trivial extension of an existing algorithm.